

COFOUNDER CLARITY PAPER

Connector

Entry Strategy & Pilot Plan (2026)

How we enter the market, what we sell first, and why customers pay



GlobalSushrut, Inc.

March 26, 2026

This document is intentionally business-first and simple. It is designed to align a cofounding team around entry wedge, pilot design, and go-to-market execution.

Contents

1	Opening	1
1.1	What this paper is trying to answer	1
1.2	What this paper is not trying to do	1
1.3	Simple framing	1
2	The Problem We Are Entering Through	3
2.1	What teams cannot answer today	3
2.2	Why this is a good entry point	4
3	Our Entry Wedge	5
3.1	Where this wedge shows up	5
3.2	What we sell	5
3.3	What we do not sell	6
4	Why This Wedge Exists Now	7
4.1	AI is moving from content generation to workflow execution	7
4.2	Risk rises as autonomy rises	7
4.3	Governance pressure is real, not theoretical	7
4.4	Why this matters for Connector	8
5	Who We Target First	9
5.1	Ideal customer profile	9
5.2	Segment 1 — AI support startups and support automation teams	9
5.3	Segment 2 — Fintech ops teams	10
5.4	Segment 3 — Internal automation teams	10
5.5	Buyer mindset	10
6	Why They Pay	12
6.1	What triggers a buying decision	12
6.2	The real buying logic	12
6.3	ROI logic they use internally	13
6.4	Why a pilot is easier to approve than a platform sale	13
7	Pilot Design	14
7.1	Core design principle — one workflow	14
7.2	What we show in the pilot	14
7.3	Pilot timeline	15
7.4	What the pilot is not	15
7.5	What success looks like	15
7.6	Pricing model for pilots	15
8	Go-To-Market Approach	17
8.1	Motion	17
8.2	How the motion works	17
8.2.1	Step 1 — Identify and reach the right teams	17
8.2.2	Step 2 — Sell one workflow, not the full vision	17

8.2.3	Step 3 — Run the pilot and land the evidence	18
8.2.4	Step 4 — Expand from pilot to commercial relationship	18
8.3	Why founder-led is correct right now	18
8.4	Phase 1 targets for 2026	18
9	Why This Works	19
9.1	We sit inside execution	19
9.2	We enforce, not just observe	19
9.3	We generate proof, not just logs	19
9.4	We start narrow and expand	19
9.5	The market timing is correct	20
9.6	What we are not doing (yet)	20
10	Collaboration and Next Step	21
10.1	Where a cofounder has immediate leverage	21
10.1.1	Shaping the wedge under market pressure	21
10.1.2	Identifying pilot customers	21
10.1.3	Refining GTM and positioning	21
10.1.4	Compressing the learning cycle	22
10.2	Roles	22
10.3	The next step	22
10.4	What this plan is not claiming	23
10.5	Closing	23

Chapter 1

Opening

This document is meant to do one job well: create alignment on where Connector enters the market, what we sell first, and how we turn early customer interest into paid pilots.

It is intentionally simple. It is not a pitch deck. It is not a technical architecture paper. It is not a full company plan.

It is a cofounder clarity document.

Core statement

We are not taking the full system to market. We are entering through a narrow wedge where AI decisions need to be traceable and provable.

Connector = execution control layer for AI decisions.

1.1 What this paper is trying to answer

After reading this, a cofounder should be able to answer three simple questions:

- Who do we call first?
- What do we sell first?
- Why will they pay?

If those answers are clear, we have something usable. If those answers are still fuzzy, the strategy is not ready.

1.2 What this paper is not trying to do

This paper is not trying to explain every part of Connector. It is not trying to carry the full technical depth of the platform. And it is not trying to impress anyone with complexity.

We already have deeper material for that. What we need here is a document that makes action easier.

Discipline point

If the reader finishes this paper understanding the architecture but still cannot explain the entry wedge, then the paper failed.

1.3 Simple framing

The market does not buy big visions first. It buys specific relief for a specific pain.

For Connector, that pain is straightforward: when an AI workflow takes action and something goes wrong, most teams cannot explain what happened, why it happened, or what evidence they can show.

That is where we enter.

Chapter 2

The Problem We Are Entering Through

AI is moving from chat to action.

It is no longer just answering questions or drafting text. It is now being used to:

- issue refunds,
- route support cases,
- approve or block transactions,
- update records,
- trigger internal workflows,
- make recommendations that directly affect customers or money.

That changes the problem.

When AI only drafts text, mistakes are mostly annoying. When AI starts taking action, mistakes become operational. They affect revenue, trust, escalations, internal confidence, and sometimes compliance.

2.1 What teams cannot answer today

When something goes wrong, most teams struggle to answer three basic questions:

1. What happened?
2. Why did it happen?
3. Can we prove it?

This sounds simple, but it is where the pain lives.

A support team cannot explain why the AI issued a refund. A fintech team cannot explain why a transaction was flagged. An internal automation team cannot show why a workflow changed path.

The result is predictable:

- automation slows down,
- people get pulled into incident review,
- buyers lose trust,
- security and compliance teams become blockers,
- AI projects lose momentum.

Human version of the problem

When something goes wrong, teams cannot clearly explain or prove what the AI did.

2.2 Why this is a good entry point

This is a good market entry problem because it is:

- real,
- already happening,
- easy to understand,
- close to business pain,
- close to budget.

We do not need to educate the market about a future problem. We need to help teams solve a present one.

That is a much better place to start.

Chapter 3

Our Entry Wedge

The wedge is the most important part of this document.

We are not entering the market as a broad AI platform. We are entering through a narrow and concrete use case:

Entry wedge

AI-driven workflows where decisions matter.

This means workflows where AI is not just speaking, but deciding or acting in ways that affect money, customers, approvals, or operations.

3.1 Where this wedge shows up

The clearest early examples are:

- support automation,
- fintech ops,
- internal copilots and internal automation.

These are good entry points because the decision has consequences. Someone feels the pain when the AI gets it wrong. And the value of a better control layer can be understood quickly.

3.2 What we sell

We sell:

What we sell first

decision trace + explanation + compliance-ready output

In practical terms, that means for one workflow we help the customer answer:

- what the AI did,
- why it did it,
- what path or context led to that output,
- what evidence can be shown to internal or external stakeholders.

This is the story in buyer language. We are not leading with models, orchestration, protocols, or internal architecture. We are leading with clarity and proof around one risky workflow.

3.3 What we do not sell

We are not selling:

- a full platform replacement,
- all AI infrastructure,
- every engine we have built,
- a giant transformation program,
- a heavy integration project.

This is important because the fastest way to lose momentum is to make Connector feel too large.

What the first conversation is not

Do not lead with: “Here is the whole Connector platform.”

Lead with: “Pick one workflow where your AI is making decisions. We will make it traceable, explainable, and provable.”

That is easier to understand, easier to buy, and easier to champion internally.

Chapter 4

Why This Wedge Exists Now

This wedge exists now because three things are changing at the same time.

4.1 AI is moving from content generation to workflow execution

Companies are no longer using AI only for drafting or summarization. They are putting it inside workflows that touch real operations.

That means the question is no longer just “Does the output look good?” The question becomes “Can we trust this inside a real process?”

External market research in 2025 and 2026 points in the same direction: AI adoption is moving deeper into business workflows, and agentic systems are becoming part of how companies think about operating leverage.

4.2 Risk rises as autonomy rises

As AI becomes more capable, weak control becomes more expensive. This is not only a security point. It is an operations point.

The more the system can do, the more a company needs:

- oversight,
- accountability,
- clear ownership,
- repeatable controls,
- evidence when something goes wrong.

PwC’s 2025 Responsible AI research is useful here because it shows governance moving closer to the teams that build and run AI. That is an important signal. Governance is not staying a policy deck. It is moving into operations.

4.3 Governance pressure is real, not theoretical

The European Commission’s AI Act timeline makes the pressure concrete. The Act is already in force, with obligations applying in phases and full applicability scheduled for August 2026.

Even outside Europe, the practical effect is larger than the law itself:

- customers ask harder questions,
- legal teams become more involved,

- compliance becomes part of vendor review,
- proof matters more than good intentions.

Simple market logic

Output is not enough. Proof is required.

4.4 Why this matters for Connector

Connector fits this moment because it is not just helping AI produce outcomes. It is helping companies understand, defend, and govern those outcomes.

That is why the wedge is timely. It sits exactly where adoption and accountability now meet.

Chapter 5

Who We Target First

We are not targeting “everyone using AI.” We are targeting teams that already feel the pain of unexplained AI decisions.

The key filter is not company size. The key filter is urgency.

5.1 Ideal customer profile

A good first customer has all three of the following:

1. AI is already live in a real workflow.
2. The workflow affects money, customers, or trust.
3. Something has already gone wrong, or the team can see that it will.

Three-question qualifier

Ask these three questions before investing time in a prospect:

1. Is your AI already taking actions, not just producing text?
2. When it makes the wrong call, can you explain why?
3. Have you had an incident yet, or do you expect one?

If the answer is yes, yes (they cannot), and yes — that is our customer.

5.2 Segment 1 — AI support startups and support automation teams

This is the clearest early wedge.

The workflow is easy to understand, the consequences are visible, and a bad AI decision shows up fast.

Real pain examples:

- a refund is issued incorrectly and the team cannot explain why,
- a discount was applied too aggressively without a clear decision trace,
- a customer escalation cannot be reconstructed from logs,
- a case was closed or routed the wrong way and the team has no trail.

Who pays: typically the VP of Support, Head of CX, or a technically strong CTO.

Why they pay: they are losing time in incident review, losing customer trust, and in regulated industries they need defensible output.

Sales cycle: short. If the pain is real and the workflow is live, a pilot conversation can open within days.

5.3 Segment 2 — Fintech ops teams

These buyers care because AI decisions in their world touch money, risk, and regulatory surface.

Real pain examples:

- transaction review where AI flags or passes decisions without a clear trace,
- fraud ops where decisions need to be defensible to internal or external auditors,
- risk scoring workflows where model behavior has to be explainable,
- exception handling where a human needs to understand why the AI routed a case.

Who pays: Head of Ops, Head of Risk, or a fintech founder in a startup context.

Why they pay: compliance and audit risk. The cost of one unexplained AI decision in a regulated fintech is higher than the cost of the tool.

Sales cycle: medium. There may be a procurement or security review, but the pain is real enough to move if we frame it correctly.

5.4 Segment 3 — Internal automation teams

These are teams using AI inside the company for approvals, case routing, exception handling, and operations.

This segment is good only when the workflow already matters and has already caused friction.

Real pain examples:

- AI-powered procurement routing with no trace when a case goes wrong,
- claims or exception handling where humans are pulled in to explain AI output,
- internal approvals driven by AI with no audit log available.

Who pays: the department head or an internal automation engineering team.

Why they pay: internal trust. Other teams will not rely on an automation they cannot understand or verify.

5.5 Buyer mindset

These buyers are **not** shopping for AI infrastructure. That language is too abstract.

They are shopping for:

- control,
- visibility,

- safety,
- explanation,
- proof.

How to position to the buyer

Do not say: “Connector is a governed agentic runtime.”

Say: “When your AI makes a decision, Connector shows what happened, why it happened, and gives you evidence you can stand behind.”

That framing is concrete, relatable, and close to a buying decision.

Chapter 6

Why They Pay

A team pays for a pilot when the cost of not solving the problem is already real.

This is the most important thing to hold in mind when having a sales conversation. They are not paying us to explore a new idea. They are paying because the pain is already expensive and they want it reduced.

6.1 What triggers a buying decision

One or more of these things is already happening for a team that will pay:

- the AI is making decisions they cannot explain,
- incidents take too long to review because the trace is not clear,
- customers are escalating AI mistakes and the team cannot respond confidently,
- sales deals are getting blocked by trust or compliance questions from prospects,
- internal teams have slowed down AI deployment because they do not feel safe,
- a compliance or legal review is approaching and there is no audit output ready.

Any one of these is a valid trigger.

6.2 The real buying logic

They pay because they want to reduce risk.

More specifically, they want to reduce:

- operational risk from unexplained AI decisions,
- customer trust risk when something goes wrong,
- internal approval friction when teams need sign-off to ship AI features,
- compliance and review friction when auditors or legal ask for evidence,
- time lost debugging AI behavior after the fact.

Core buying signal

They are not buying a tool.

They are buying a reduction in a specific, visible, already-expensive risk.

6.3 ROI logic they use internally

When a buyer takes this to their team for approval, they frame it one of these ways:

1. **Incident cost reduction.** One reviewed incident costs us X hours. If this reduces reviews by Y per month, the saving is clear.
2. **Compliance enablement.** We need audit-ready output for our next review. This gives us that at lower cost than building it.
3. **Sales unblocking.** We lost two deals because prospects asked for AI decision traceability and we had nothing to show. This solves it.
4. **Confidence to ship faster.** Our team will not deploy more AI workflows without a control layer. This removes the blocker.

Any one of these frames is sufficient. The buyer does not need all four.

6.4 Why a pilot is easier to approve than a platform sale

A platform sale sounds large, uncertain, and expensive. A single-workflow pilot sounds:

- bounded,
- measurable,
- low risk,
- useful immediately.

That is the exact shape of a budget approval that gets through quickly.

Pilot framing that works

Do not say: “Evaluate our platform.”

Say: “Let us take one workflow where your AI is making decisions and make it traceable and explainable in two to four weeks.”

That is a much easier yes.

Chapter 7

Pilot Design

The pilot is the most important commercial moment in 2026. Getting the pilot structure right matters more than getting the product perfect.

The goal of the pilot is not to show everything Connector can do. The goal is to prove one thing clearly:

Pilot goal

After this pilot, the customer can see what their AI did and why.

7.1 Core design principle — one workflow

The right pilot is:

- one workflow,
- one team,
- one operational problem,
- one success definition.

Not multi-team. Not a broad rollout. Not “let us instrument the whole company.”

The narrower the pilot, the faster it succeeds. The faster it succeeds, the easier expansion becomes.

7.2 What we show in the pilot

For the chosen workflow, we show:

1. **Full trace.** Every step the AI took, in order, with data context.
2. **Explanation of the decision path.** Why the AI reached the output it reached.
3. **Audit-ready output.** Something a non-technical stakeholder can review and understand.
4. **Clearer operational visibility** than the team has today without Connector.

The customer should come away with a simple reaction:

The reaction we want

“Now we can see what happened.”

That reaction is more valuable than any technical feature demonstration.

7.3 Pilot timeline

Context	Timeline
Fast-moving startup or internal team	2–3 weeks
Mid-size company with some review process	3–4 weeks
Larger org with security or procurement review	4–6 weeks

The principle is the same regardless of timeline: short, narrow, concrete.

7.4 What the pilot is not

- It is not a proof of concept for the whole platform.
- It is not a free engagement.
- It is not indefinite.
- It is not a replacement for a commercial agreement.

The pilot has a defined scope, a defined timeline, and a defined success criteria.

7.5 What success looks like

A pilot is complete and successful when the customer can say:

1. We can reconstruct what the AI did inside this workflow.
2. We can explain why it reached that output.
3. We can show this evidence to internal stakeholders.
4. We want this on at least one more workflow.

The fourth point is the expansion bridge. That is where a pilot becomes a commercial relationship.

7.6 Pricing model for pilots

The pilot should be paid. A free pilot signals low confidence and attracts low-quality evaluation.

A reasonable pilot price is between \$2,000 and \$8,000 depending on the segment:

Segment	Pilot price range
AI support startup	\$2,000–\$4,000
Fintech ops team	\$4,000–\$8,000
Internal automation at a mid-market company	\$3,000–\$6,000

A paid pilot also creates a clear transition path to a monthly or annual contract if it succeeds.

Chapter 8

Go-To-Market Approach

The motion should be simple and realistic. We are not running a broad campaign. We are not waiting for the market to discover us. We are going directly to the right people with the right message.

8.1 Motion

- Founder-led.
- Direct outreach.
- Network-based.

This is the right motion because the wedge is still being validated, the message is still being refined, and founder conversations generate the fastest learning at this stage.

8.2 How the motion works

8.2.1 Step 1 — Identify and reach the right teams

We target teams that already fit the wedge profile from Chapter 5.

Outreach should be direct and specific. Not “we built a powerful AI platform.”

Instead:

Outreach framing that works

“Are you using AI to take actions in support or ops?
When it makes the wrong call, can you explain why?
If not, we should talk.”

This framing pre-qualifies the conversation and makes the pain immediately visible to the right buyer.

8.2.2 Step 2 — Sell one workflow, not the full vision

The first sale is:

- one workflow,
- one pain point,
- one pilot.

This keeps conversations manageable, keeps the buyer’s internal champion in control, and keeps us focused on proving value rather than educating.

Selling the full platform vision in the first conversation almost always slows things down.

8.2.3 Step 3 — Run the pilot and land the evidence

After the pilot succeeds:

- document what was shown,
- get a clear statement from the customer,
- use it in the next conversation.

One strong pilot outcome is worth more than ten vague conversations.

8.2.4 Step 4 — Expand from pilot to commercial relationship

After a successful pilot, expansion happens in one of three ways:

1. More workflows in the same team.
2. More teams in the same company.
3. A reference story that opens doors at similar companies.

This is where Connector becomes sticky. Once a company relies on us to make AI decisions explainable and defensible, we become part of how they run AI safely.

8.3 Why founder-led is correct right now

At this stage, the founder should be in most early conversations because:

- the wedge language is still being sharpened,
- objections reveal what needs to change in positioning,
- the pilot design benefits from direct customer input,
- early buyers often want to know who is building this.

Founder-led is not a weakness. It is the right instrument for the validation phase.

8.4 Phase 1 targets for 2026

By end of Q2 2026:

- 3–5 qualified conversations with teams in the target segments,
- 1–2 paid pilots in progress or completed,
- at least one clear pilot win documented with a customer statement.

By end of Q3 2026:

- 2–3 pilots completed with documented outcomes,
- at least 1 expansion conversation underway,
- a tightened message based on what objections surfaced most often.

These are not aggressive targets. They are honest targets for a founder-led wedge validation phase.

Chapter 9

Why This Works

There are a small number of reasons why this approach works. Keeping them simple helps us stay focused.

9.1 We sit inside execution

Connector is not a dashboard bolted onto the outside of an AI workflow. It sits close to the moment where decisions happen.

That matters because traceability that is imposed after the fact is fragile. Evidence that is generated at the point of execution is reliable.

A buyer who wants defensible AI decision output needs something that was present when the decision happened. Connector is that layer.

9.2 We enforce, not just observe

Most observability and monitoring tools watch what AI does. Connector does more than watch.

It operates at the control layer. That means it can enforce policies, capture evidence, and produce proof in real time rather than reconstructing it afterward from incomplete logs.

That is a stronger value proposition for compliance and audit.

9.3 We generate proof, not just logs

Logs are traces. Proof is traceable, signed, structured output that a non-technical stakeholder can review.

There is a large gap between those two things.

When a customer needs to respond to a customer complaint, a compliance review, or an internal escalation, they need proof, not logs. Connector produces the right output for that situation.

9.4 We start narrow and expand

The pilot model means we never try to win the whole account on the first conversation. We win one workflow. Then we expand.

This is how infrastructure tools build durable revenue: they land small, prove value fast, and expand as trust grows.

9.5 The market timing is correct

As Chapter 4 explains, three shifts are happening at the same time:

1. AI is moving into workflows where decisions have consequences.
2. Governance pressure is becoming real, not just theoretical.
3. Buyers are learning that output quality is not sufficient.

Connector's wedge sits exactly at the intersection of those three shifts.

Why Connector wins in the wedge

We are present at execution.

We produce structured, signed evidence.

We start on one workflow and expand.

The market timing is right.

9.6 What we are not doing (yet)

This section matters because it shows focus.

We are **not** doing:

- full-platform GTM,
- all-industry targeting,
- broad enterprise rollout from day one,
- aggressive scaling before wedge validation,
- complex multi-product positioning in early conversations.

We are also not asking the buyer to understand the full architecture. We are asking them to let us solve one specific problem in one specific workflow.

That is a much smaller ask. And it is the right ask at this stage.

Chapter 10

Collaboration and Next Step

10.1 Where a cofounder has immediate leverage

The value of a cofounder at this stage is not just vision alignment. It is making the wedge commercially real.

There are four areas where the contribution matters most right now.

10.1.1 Shaping the wedge under market pressure

The wedge is a hypothesis until a customer pays. A cofounder who is willing to have direct conversations helps pressure-test:

- which workflows are most painful,
- which segments move fastest,
- which version of the story gets meetings,
- which version of the pitch reduces objections.

10.1.2 Identifying pilot customers

The most valuable short-term contribution is identifying two to three credible pilot opportunities where:

- AI is already live in a real workflow,
- the decision has consequences,
- the buyer can move within weeks,
- the workflow can be scoped tightly.

This matters more than broad lead generation. Quality of the first conversations is more important than volume.

10.1.3 Refining GTM and positioning

The practical GTM questions need answers:

- Who do we call first?
- Which segment lands fastest?
- What wording creates immediate recognition of the pain?
- What pilot package is easiest to approve?

- What outcome do we promise and how do we measure it?

These are not strategic questions to revisit quarterly. These are operational questions to answer in weeks and refine continuously.

10.1.4 Compressing the learning cycle

We need fast feedback loops from real conversations. A cofounder accelerates this because:

- two people can have more conversations simultaneously,
- two perspectives on the same conversation improve interpretation,
- faster iteration on the pitch means faster clarity on what works.

10.2 Roles

Area	Founder	Cofounder
Product and architecture	Primary	Advisory
Pilot design and scoping	Primary	Supporting
Outreach and pipeline	Shared	Shared
Positioning and messaging	Shared	Shared
Pilot customer identification	Supporting	Primary
Network activation	Supporting	Primary

10.3 The next step

The next step is narrow and concrete.

Next step

Identify two to three pilot opportunities and validate the wedge in real workflows.

Specifically, we need to confirm:

1. **Who we call first.** Which teams in which segments are most likely to feel the pain right now?
2. **What we sell first.** Which workflow framing gets the fastest yes?
3. **Why they will pay.** Which version of the risk-reduction story closes the conversation?

If we can answer those three questions with real customer conversations and one or two paid pilots, the rest of the commercial path becomes much clearer.

10.4 What this plan is not claiming

This plan is not claiming that every assumption here is correct. Some segments may move slower than expected. Some framing may need adjustment. The pilot timeline may vary.

That is fine. The purpose of this plan is to create enough clarity to take action.

Once we are in real conversations, we update based on what we learn. That is how an early market entry actually works.

10.5 Closing

The purpose of this document was simple:

- Explain where Connector enters.
- Explain what we sell.
- Explain why someone pays.
- Explain what a pilot looks like.
- Explain how we go to market.
- Explain where a cofounder plugs in.

If those six things are clear after reading this, the document did its job.

One-line summary

“Pick one AI workflow where decisions matter. We make it traceable, explainable, and provable. Then we expand.”

That is the whole strategy. Everything else is execution.