

Connector

The operating layer for intelligence that already acts

Buyer brief — August 2026. What the system can do, why they need it, who buys, what we sell, why us.

1. What we are selling

A self-hosted governed intelligence kernel.

Customers keep LangGraph, Crew, vLLM, Ollama, OpenAI, MCP. Those products do what they do. Connector sits **beneath** them and does the job they skip:

- **Charter** an intelligence I (identity μ , not a Linux PID).
- **Admit** every effect A (talk, tool, world, share) — or refuse it.
- **Keep** the world model ($/m$, $/k$) that outlives one graph.
- **Stop** this I (interrupt thought, kill-switch, compensating undo).
- **Prove** who did what (hash-chained traces; council floor with live μ).

SKU: one node (`connector-platform` + `connectorctl`), licensed, self-hosted. Not a chatbot. Not a GPU farm. Not Honor/CosmOS. Not “we are LangGraph.”

Claim language we may use: governed intelligence kernel (V1 buyer Agent OS). We must not say: Mei et al. AIOS, LLM-as-CPU, court-green, SOC2 certified, Honor Agentic OS, infrastructure GPU OS.

2. What the system can do today

These are in code. Crew and AutoGen put them on slides.

Capability	What that means in production	How they run it without us
Kernel who-am-I	Talk injects identity μ (0xCD). An I cannot speak as another.	A persona prompt. The model invents a name.
Isolation by default	Agent A cannot see Agent B. Share only after human+root portal.	One Python process. Shared graph RAM.
World grants	Each ($I \times$ address) is its own grant. MCP behind admit.	Ambient tools on the graph.
Gartner L1–L4	Root/Cone = L3 (act with approval). App = L4 with kill + compensate. Per crossing, not a global switch.	Uniform lock-or-trust. Gartner said that fails.
Stop this I	Kill-switch interrupts this intelligence’s thought. Seconds on this node.	Restart the process. Everyone dies.
Compensate	Revoke grant, close portal, deny tool. WM records the line. Not world rewind.	They say “rollback.” They do not rewind the world.
Cost cap	Per- I budget gate on Talk.	A dashboard after the bill.
Auditor trail	DecisionTrace + OTel/JSONL export. Council floor is who-said-what with μ .	“The framework logged it somewhere.”
Memory that lasts	$/m$ core, recall, archival; Talk inject; kernel paging.	RAG stuffed by the app. Dies with the graph.
Council	Root-minted members. Named task owner μ . Talk injects the desk.	Crew chat loop. Invented names.
Sit beneath	<code>OPENAI_BASE_URL=/v1</code> . Swap vLLM $\leftrightarrow$ Ollama $\leftrightarrow$ OpenAI without changing I .	Lock-in to their runtime.

Operator loop: charter $I \rightarrow$ link device $\rightarrow$ point the app at `/v1` $\rightarrow$ grant world $\rightarrow$ admit + trace $\rightarrow$ interrupt or kill this I . Follow: `connectorctl iia aios · GET /api/v1/kernel/aios/claim-readiness` $\rightarrow$ `buyer_surface`.

3. Why they need it

Gartner (25 Jun 2025): **over 40% of agentic projects canceled by end-2027** — cost, unclear value, or inadequate risk controls. Gartner (26 May 2026): applying **uniform** governance across agents will fail; map each agent to Observe / Advise / Act-with-approval / Act-autonomous; L4 needs rollback, circuit breaker, named owner, and a **five-minute stop**. The buyer already has agents that **act**: they spend, read data, call tools, change customer state. The framework built the

workflow. It did not build:

- 1. a stop that is this agent, not the whole box;
- 2. a cost cap before the invoice;
- 3. isolation so two agents do not share a brain;
- 4. an identity the model cannot invent;
- 5. a trail an auditor will sit with (not a PDF policy, not a Slack export).

They need Connector because **the coordination layer is the remaining lock-in** once MCP made tools portable. Without it they cannot ship L4, cannot answer “who did that,” and join the 40%.

4. Who will buy

Not “everyone using ChatGPT.” The filter is **urgency**: AI already takes action, the action can hurt, and they cannot explain the last miss.

Buyer	Why they feel it	Who signs
Support / CX automation	Wrong refund, wrong discount, unreconstructable escalation.	VP Support, Head of CX, CTO
Fintech / ops risk	Flags, exceptions, money movement without a trail.	Head of Ops, CRO, CISO
Platform / AI engineering	Many agents, three frameworks, no OS. Security will not sign L4.	VP Eng, Head of Platform
Compliance / legal (influencer)	EU AI Act inventory, evidence pack, “show me who approved.” They do not write the PO; they can kill the deal.	GC, CISO, DPO

Three-question qualifier. (1) Is the AI taking actions, not only producing text? (2) When it is wrong, can you say which intelligence did it? (3) Have you had an incident, or can you see the one coming? Yes / no / yes — that is a customer.

Not first: teams still in chatbot PoC; buyers who want Honor-style device UX; buyers who want us to become their GPU/vLLM farm.

5. Why us

- **We are the operating layer, not a fourth framework.** LangGraph and Crew stay. Identity, WM, admit, and stop do not live in their graph.
- **CPU, GPU, and intelligence are three kinds.** We do not sell “LLM as CPU.” vLLM is a device URL. *I* is chartered intelligence. That is why we can swap engines without rewriting the fleet.
- **Isolation is default, not a prompt.** A sharing contract (human+root) mints a pore. Council is the same rule: named owner μ , not a shared crew RAM.
- **Evidence is fail-closed.** Court-readiness never auto-greens. Art.9 inventory is a living list, not a spreadsheet. We will not sell a green checkbox as a certificate.
- **It is already running.** The bake-off list is **buyer_surface** in the claim JSON — not a roadmap slide.

6. What we do not sell

- A rewrite of the customer’s LangGraph / Crew / vLLM stack.
- World rewind (undo is compensating).
- Certified SOC2 / ISO / court-green from this brief.
- Honor / CosmOS device OS. Infrastructure GPU OS.
- Live multi-node fabric as the default (product SoT is single-node until soak + flag).

7. How they buy

Entry: one live workflow where AI already acts and a miss is expensive. We put that workflow on three sockets (completion /v1, syscall, world grant). They keep their thinker. They get identity, admit, stop, budget, and a trail.

Pilot shape: 30 days, self-hosted node, one *I* (then two, to prove isolation and council). Success = they can stop this *I*, see who did what, and show a crossing that was refused.

Expansion: more *I*, more grants, council, SCIM/OIDC, Art.9 inventory as packaging — same kernel.

