

Connector Design-Partner Target Universe

100 North American Companies Matching ICP1, ICP2, and ICP3

USA and Canada only

How to read this report

This report is built for founder-led outbound and design-partner selection. It follows the ICP definitions in the internal workflow control deck:

- **ICP1:** companies already running internal AI workflows or agentic systems in production and actively fixing instability.
- **ICP2:** experimentation-heavy AI labs and platform companies building agents, evals, or workflow infrastructure where internal iteration speed and reliability are core.
- **ICP3:** regulated healthcare, fintech, and legal companies where internal AI workflows need traceability, policy control, and explainability.

The fit call is directional and based on public signals from company product pages, launch posts, case studies, YC or venture directories, and public positioning. It is meant for prospecting and prioritization, not certainty.

Difficulty legend

- **Access:** how hard it is to get a serious design-partner conversation with the right buyer.
- **Onboarding:** how hard it is to get technically onboarded into a real workflow fast enough to prove value.
- **Low:** approachable team, narrower scope, simpler integration.
- **Medium:** credible target, but requires clear use case, buyer access, and workflow mapping.
- **High:** enterprise buyer motion, complex integrations, or sensitive systems.
- **Very High:** large incumbent, strong internal platform teams, or regulated implementation burden.

Executive summary

- **ICP1 total:** 50 companies.
- **ICP2 total:** 30 companies.
- **ICP3 total:** 20 companies.
- **Best early design-partner profile:** mid-market to growth-stage platform companies already shipping AI agents into customer support, internal work AI, or workflow automation, where the team already understands observability but still lacks runtime control.
- **Highest-probability early wedge:** support automation, contact center AI, enterprise work AI, and agent orchestration vendors with visible workflow complexity and public language around safeguards, evals, or governance.

ICP1: Internal Fix Loop — 50 companies

Ada

- **HQ / ICP:** Toronto, Canada / ICP1
- **Why it matches:** Ada is already in enterprise customer service automation with autonomous decisioning, which means real production workflow reliability is core to the product.
- **Public signal:** Public product pages emphasize AI agents, autonomous decision making, adaptive reasoning, and multi-layer safeguards.
- **Likely buyer:** VP Engineering, Head of AI, CTO, VP CX Technology
- **Difficulty:** Access: Medium; Onboarding: Medium-High
- **Why difficulty is at that level:** The category fit is very direct, but proving value will require access to live support flows, escalation rules, and policy boundaries.

Glean

- **HQ / ICP:** Palo Alto, USA / ICP1
- **Why it matches:** Glean is explicitly building enterprise agents that operate across business systems, so orchestration and execution control are central problems.
- **Public signal:** Glean publicly markets building, deploying, orchestrating, and governing AI agents inside enterprise workflows.
- **Likely buyer:** Head of AI Platform, VP Engineering, CTO
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** Strong fit, but security, internal knowledge graph access, and cross-system execution make evaluation and rollout heavier.

Moveworks

- **HQ / ICP:** Mountain View, USA / ICP1
- **Why it matches:** Moveworks already runs AI across internal employee workflows such as IT and HR, where bad actions and brittle execution are expensive.
- **Public signal:** Public messaging describes a reasoning engine that understands, plans, executes, and adapts inside enterprise workflows.
- **Likely buyer:** VP Product, Head of AI Platform, CTO
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** Buyer quality is high, but the company already has serious platform depth and a complex integration environment.

Aisera

- **HQ / ICP:** Palo Alto, USA / ICP1
- **Why it matches:** Aisera is a close fit because it automates multi-step workflows across IT, HR, finance, and support.
- **Public signal:** Its public platform pages highlight AI workflows, agentic reasoning, orchestration, and hyperflow catalogs.
- **Likely buyer:** VP Engineering, CIO, Head of Automation
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** The fit is strong, but deployments usually span multiple systems and require enterprise security review.

Sierra

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Sierra is building customer experience agents with explicit decisioning and workflow logic, which is where runtime control matters.
- **Public signal:** Sierra describes decisioning engines, next-best-action workflows, and support for engineering teams building agents.
- **Likely buyer:** VP Engineering, CTO, Head of CX Platform
- **Difficulty:** Access: Medium-High; Onboarding: Medium-High
- **Why difficulty is at that level:** Still design-partner plausible, but the company is well-capitalized and technically mature enough to expect a sharp wedge.

Decagon

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Decagon is one of the clearest fits because it already speaks the language of reliability, testing, observability, and iteration for AI agents.
- **Public signal:** Public materials emphasize robust testing, observability, experimentation, and agent reliability as systems evolve.
- **Likely buyer:** Head of AI Engineering, VP Product, CTO
- **Difficulty:** Access: Medium; Onboarding: Medium-High
- **Why difficulty is at that level:** The conceptual fit is excellent, but the bar for technical novelty will be high because they already understand the problem deeply.

Coveo

- **HQ / ICP:** Quebec, Canada / ICP1
- **Why it matches:** Coveo sits near agentic search, self-service, and agent workflows, where retrieval quality and action quality intersect.
- **Public signal:** Public pages mention agentic AI, self-service, agent workflows, and sovereign AI partnerships for regulated buyers.
- **Likely buyer:** VP Product, CTO, Head of AI Relevance
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** The buyer case is real, but integrations into enterprise relevance and workflow layers will be non-trivial.

Forethought

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Forethought is a strong fit because workflow duplication, wrong answers, and fact validation are already visible pain points in the category.
- **Public signal:** Public platform materials talk about multi-agent systems, verification of facts, fallback workflows, and support automation complexity.
- **Likely buyer:** VP Support Technology, CTO, GM AI
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** The fit is direct and pain is public, but the scope still requires access to live support routing and automation logic.

Netomi

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Netomi is already positioning around policy-aware agentic customer service, which is close to the runtime control wedge.
- **Public signal:** Public messaging highlights agentic AI, policies, guardrails, proactive action, and customer experience automation.
- **Likely buyer:** VP CX Operations, CTO, Head of AI Platform
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** A credible target with real workflow execution, though a first pilot still needs a tightly bounded use case.

Cresta

- **HQ / ICP:** Palo Alto, USA / ICP1
- **Why it matches:** Cresta blends human and AI agents in live contact center operations where workflow reliability is commercially meaningful.
- **Public signal:** Public pages highlight workflow automation, quality management automation, and AI for customer conversations.
- **Likely buyer:** VP Contact Center Technology, CTO, VP Operations
- **Difficulty:** Access: Medium; Onboarding: Medium-High
- **Why difficulty is at that level:** Buyer pain is clear, but production contact center stacks tend to be integration-heavy and KPI sensitive.

Observe.AI

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Observe.AI is an especially good fit because it talks about predictable outcomes, mission-critical steps, and workflow enforcement in contact centers.
- **Public signal:** Public platform copy says its agents capture intent accurately, enforce mission-critical steps, integrate broadly, and improve continuously.
- **Likely buyer:** VP CX Operations, Head of AI Platform, CTO
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** The fit is strong, but real proof will require deep access to live workflows and operational metrics.

Gorgias

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Gorgias runs ecommerce support and sales automation where bad workflow execution directly hits customer experience and revenue.
- **Public signal:** Public product pages say the AI agent learns a brand, the customer, and the workflow, and automates interactions continuously.
- **Likely buyer:** VP Support, VP Engineering, Head of Automation
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** The scope is tangible and buyer pain is real, but success depends on linking into existing ecommerce and support logic.

Zendesk

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Zendesk is moving from basic automation toward self-improving AI agents, which makes execution control a natural next layer.
- **Public signal:** Public materials describe AI agents, a resolution learning loop, and optimization before deployment.
- **Likely buyer:** GM AI, VP Platform, VP Product
- **Difficulty:** Access: Very High; Onboarding: High
- **Why difficulty is at that level:** Massive installed base and resources make this a strategic target, but not an easy first design partner.

ServiceNow

- **HQ / ICP:** Santa Clara, USA / ICP1
- **Why it matches:** ServiceNow already owns enterprise workflow infrastructure, so agentic execution quality is a real and expensive concern.
- **Public signal:** Public positioning around workflow automation and AI assistance shows the company is moving deeper into agentic work.
- **Likely buyer:** SVP Platform, VP Workflow AI, CIO Office
- **Difficulty:** Access: Very High; Onboarding: Very High
- **Why difficulty is at that level:** The fit is real, but the company is large, process-heavy, and likely to expect a very polished product story.

Salesforce

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Salesforce is pushing Agentforce into production CRM workflows, which are highly valuable and failure-sensitive.
- **Public signal:** Public messaging focuses on AI agents operating across CRM processes and customer operations.
- **Likely buyer:** SVP AI Platform, VP Product, CIO Office
- **Difficulty:** Access: Very High; Onboarding: High
- **Why difficulty is at that level:** A marquee target, but the internal platform depth and buyer sophistication make early entry hard.

Freshworks

- **HQ / ICP:** San Mateo, USA / ICP1
- **Why it matches:** Freshworks has AI deeply embedded into customer and employee service workflows, which creates runtime reliability needs.
- **Public signal:** Its public AI positioning emphasizes agentic automation across support and service desks.
- **Likely buyer:** VP Product, CTO, Head of AI
- **Difficulty:** Access: Medium-High; Onboarding: Medium
- **Why difficulty is at that level:** Good category fit with less institutional complexity than the biggest incumbents.

Kore.ai

- **HQ / ICP:** Orlando, USA / ICP1
- **Why it matches:** Kore.ai operates in enterprise virtual assistants and automation, where orchestration and reliability are everyday problems.
- **Public signal:** Public positioning centers on enterprise AI agents and workflow execution across service functions.
- **Likely buyer:** Head of Conversational AI, CTO, VP Product
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** Fit is direct, but platform breadth and integration scope raise onboarding complexity.

Amelia

- **HQ / ICP:** New York, USA / ICP1
- **Why it matches:** Amelia is in digital employees and enterprise automation, both of which need better runtime decision control.
- **Public signal:** Public pages highlight enterprise AI agents across service workflows and automation-heavy use cases.
- **Likely buyer:** SVP Automation, VP Engineering, CIO
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** Large workflow footprint and enterprise deployment style increase implementation burden.

Talkdesk

- **HQ / ICP:** Palo Alto, USA / ICP1
- **Why it matches:** Talkdesk sits in high-volume contact center workflows where bad execution hurts both cost and customer satisfaction.
- **Public signal:** Public messaging emphasizes AI-driven contact center automation and virtual agents.
- **Likely buyer:** VP Contact Center Technology, CTO
- **Difficulty:** Access: High; Onboarding: Medium-High
- **Why difficulty is at that level:** Good fit, but competitive noise is high and production contact center integrations are not lightweight.

Kustomer

- **HQ / ICP:** New York, USA / ICP1
- **Why it matches:** Kustomer operates customer service workflows where AI triage, routing, and action quality can be improved materially.
- **Public signal:** Public messaging centers on AI-powered customer service workflows and automation.
- **Likely buyer:** VP CX Engineering, Head of Support Systems
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** The workflow surface is clear and bounded enough for a credible pilot if the right champion exists.

Intercom

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Intercom is a strong fit because Fin AI Agent is already built to sit inside workflows, triage logic, and handoff logic.
- **Public signal:** Intercom docs explicitly describe using Fin AI Agent inside Workflows for customization and triage.
- **Likely buyer:** VP Support Engineering, CTO, Head of AI Product
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** The fit is obvious, but the bar is higher because they already understand workflow automation deeply.

Assembled

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Assembled is close to support operations and AI workforce management, where teams are actively tuning automation quality.
- **Public signal:** Public positioning combines support operations software with AI features for modern teams.
- **Likely buyer:** Head of Support Operations, VP Engineering
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** A practical design-partner target because the workflow scope is concrete and operational pain is measurable.

Retool

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Retool is compelling because it sits in internal tools and operational workflows where agent behavior meets real systems quickly.
- **Public signal:** Public positioning emphasizes building internal apps, workflows, and AI-powered operations surfaces.
- **Likely buyer:** Head of Internal Tools, VP Engineering, CTO
- **Difficulty:** Access: Medium; Onboarding: High
- **Why difficulty is at that level:** If the wedge lands, it can be strategic, but pilots will require serious integration into internal tools and permissions.

Workato

- **HQ / ICP:** Mountain View, USA / ICP1
- **Why it matches:** Workato is already an orchestration layer for enterprise automation, making it naturally adjacent to execution control.
- **Public signal:** Public materials emphasize AI agents and automation orchestration across enterprise systems.
- **Likely buyer:** VP Automation, Enterprise Architect, CTO
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** Great conceptual fit, but workflow breadth and incumbent platform depth raise the proof burden.

Zapier

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Zapier is a good fit because agentic automation is now a direct extension of its workflow platform.
- **Public signal:** Public positioning now includes AI agents and workflow automation across thousands of apps.
- **Likely buyer:** Head of AI Automation, CTO, VP Product
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** Broad workflow access is useful, but the product is already close to the problem space.

UiPath

- **HQ / ICP:** New York, USA / ICP1
- **Why it matches:** UiPath is one of the strongest conceptual fits because agentic automation sits directly on top of execution-heavy workflows.
- **Public signal:** Public messaging focuses on agentic automation and orchestration across enterprise processes.
- **Likely buyer:** VP Intelligent Automation, CIO, CTO
- **Difficulty:** Access: High; Onboarding: Very High
- **Why difficulty is at that level:** The fit is strategic, but the platform is massive and onboarding into meaningful workflows will be demanding.

Pegasystems

- **HQ / ICP:** Cambridge, USA / ICP1
- **Why it matches:** Pega is already deep in decisioning and enterprise workflow execution, so runtime AI control is a natural adjacency.
- **Public signal:** Public positioning highlights enterprise process orchestration, decisioning, and AI-enhanced automation.
- **Likely buyer:** VP Process Automation, CIO Office, CTO
- **Difficulty:** Access: High; Onboarding: Very High
- **Why difficulty is at that level:** Buyer sophistication is high and implementation environments are complex.

Gong

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Gong is moving toward revenue AI that touches live go-to-market workflows, where broken execution can be expensive.
- **Public signal:** Public positioning emphasizes revenue AI, action recommendations, and automation across sales operations.
- **Likely buyer:** VP Revenue Systems, CTO, Head of AI Product
- **Difficulty:** Access: Medium-High; Onboarding: Medium
- **Why difficulty is at that level:** The workflow value is real, but the pain must be tied to operational execution rather than analytics alone.

Clari

- **HQ / ICP:** Sunnyvale, USA / ICP1
- **Why it matches:** Clari is a fit where AI recommendations and revenue workflows move from visibility toward action.
- **Public signal:** Public messaging increasingly ties AI to execution and revenue process automation.
- **Likely buyer:** VP Revenue Operations, CTO
- **Difficulty:** Access: Medium-High; Onboarding: Medium
- **Why difficulty is at that level:** The company fits best if the pilot is framed around execution and not generic forecasting or insights.

Outreach

- **HQ / ICP:** Seattle, USA / ICP1
- **Why it matches:** Outreach runs AI-assisted sales execution workflows where action quality and sequencing matter.
- **Public signal:** Public product pages emphasize AI-driven prospecting and revenue execution.
- **Likely buyer:** VP RevOps, CTO, Head of AI Product
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** The fit is decent, but value must be shown at the workflow-control layer, not just productivity.

Salesloft

- **HQ / ICP:** Atlanta, USA / ICP1
- **Why it matches:** Salesloft is similar to Outreach: AI-guided workflows with real execution steps inside sales operations.
- **Public signal:** Public materials emphasize AI sales automation and guided execution.
- **Likely buyer:** VP RevOps, CTO
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** Reachable if framed as reliability and control in revenue workflows rather than generic AI.

HubSpot

- **HQ / ICP:** Cambridge, USA / ICP1
- **Why it matches:** HubSpot is putting AI agents and automation into GTM and support workflows, which creates runtime control needs.
- **Public signal:** Public product launches around Breeze and AI agents show growing workflow automation ambition.
- **Likely buyer:** GM AI, VP Product, Head of Platform
- **Difficulty:** Access: High; Onboarding: Medium
- **Why difficulty is at that level:** Large brand and broad surface area increase access difficulty, though onboarding can be narrowed to one flow.

Front

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Front lives in shared inbox and operations workflows where AI routing and actioning can fail quietly.
- **Public signal:** Public positioning includes AI and workflow automation for customer operations teams.
- **Likely buyer:** VP Operations Systems, Head of Support Technology
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** The target is reachable and workflow scope is concrete enough for a bounded pilot.

Rippling

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Rippling owns cross-functional workflows across HR, IT, and finance where AI mistakes can have real operational impact.
- **Public signal:** Public product direction increasingly combines automation, policy, and AI in employee operations.
- **Likely buyer:** CTO, Head of Platform, CIO Office
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** The use case is strong, but internal systems and permissions make onboarding meaningfully harder.

Paradox

- **HQ / ICP:** Scottsdale, USA / ICP1
- **Why it matches:** Paradox automates recruiting workflows with conversational AI, a strong example of internal execution loops under active tuning.
- **Public signal:** Public positioning centers on AI assistants automating hiring operations.
- **Likely buyer:** VP Talent Technology, CTO
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** Good fit with a bounded domain, though the wedge must be tied to workflow correctness not just automation.

Qualified

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Qualified uses AI in pipeline-generation workflows where wrong execution can hurt lead handling and conversion.
- **Public signal:** Public messaging emphasizes AI SDR and pipeline automation.
- **Likely buyer:** VP Growth Engineering, VP RevOps
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** The workflow surface is clear, but the design-partner angle has to be specific about runtime control.

Gladly

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Gladly is in customer service operations where AI and workflow quality are directly monetizable.
- **Public signal:** Public positioning highlights AI-driven service experience and automation.
- **Likely buyer:** VP CX Technology, CTO
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** Strong enough fit if the pilot is tied to one live service flow with measurable containment or resolution gains.

Symbl.ai

- **HQ / ICP:** Seattle, USA / ICP1
- **Why it matches:** Symbl.ai is close to workflow-aware conversation intelligence and real-time operational assistance.
- **Public signal:** Public positioning focuses on real-time conversation intelligence and workflow-oriented APIs.
- **Likely buyer:** Head of AI Platform, VP Product
- **Difficulty:** Access: Medium; Onboarding: Medium-High
- **Why difficulty is at that level:** Useful fit, but the wedge has to connect to execution control rather than raw conversation analytics.

Ema

- **HQ / ICP:** Mountain View, USA / ICP1
- **Why it matches:** Ema is building general enterprise AI employees, which makes runtime decision control a foundational problem.
- **Public signal:** Public positioning emphasizes universal AI employees operating across functions and systems.
- **Likely buyer:** CTO, Head of AI Transformation, VP Engineering
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** The promise is broad, so onboarding a meaningful pilot will require carefully narrowing the first workflow.

DevRev

- **HQ / ICP:** Palo Alto, USA / ICP1
- **Why it matches:** DevRev is an AI-native support and product operations platform where workflow execution is central.
- **Public signal:** Public positioning links conversational AI, support, and product workflows.
- **Likely buyer:** VP Support Engineering, CTO
- **Difficulty:** Access: Medium; Onboarding: Medium-High
- **Why difficulty is at that level:** Strong fit, but the company is building an AI-native stack and will expect clear technical differentiation.

Maven AGI

- **HQ / ICP:** Boston, USA / ICP1
- **Why it matches:** Maven AGI is directly in enterprise support automation, which is one of the cleanest execution-control wedges.
- **Public signal:** Public positioning focuses on agentic support automation for enterprise teams.
- **Likely buyer:** VP CX Engineering, CTO
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** Reasonably reachable and highly relevant, especially if the pilot focuses on a narrow support flow.

Meya

- **HQ / ICP:** Canada / ICP1
- **Why it matches:** Meya is a programmable support automation platform, so workflow logic and runtime behavior are first-order concerns.
- **Public signal:** Public positioning emphasizes customer service automation and flexible workflow logic.
- **Likely buyer:** Head of Support Technology, CTO
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** The fit is good and scope is controllable, but a clear pain point needs to be surfaced quickly.

Pypestream

- **HQ / ICP:** New York, USA / ICP1
- **Why it matches:** Pypestream operates conversational automation in enterprise service workflows where bad execution has visible cost.
- **Public signal:** Public positioning centers on enterprise customer engagement automation and workflow execution.
- **Likely buyer:** VP Digital Service, Head of Automation
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** Enterprise deployment style increases onboarding work even though the wedge is relevant.

LivePerson

- **HQ / ICP:** New York, USA / ICP1
- **Why it matches:** LivePerson has long been in conversational AI and customer automation, making workflow reliability a real issue.
- **Public signal:** Public positioning connects AI agents to conversational workflows at enterprise scale.
- **Likely buyer:** SVP Conversational AI, CTO
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** Large public-company motion makes access harder and implementation heavier.

Verint

- **HQ / ICP:** New York, USA / ICP1
- **Why it matches:** Verint is a credible fit through customer engagement and AI-powered automation across large service operations.
- **Public signal:** Public positioning highlights CX automation, AI bots, and workflow optimization.

- **Likely buyer:** VP Contact Center Platform, CIO Office
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** The company is large and enterprise-focused, so pilots will need strong executive sponsorship.

Fullview

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Fullview is closer to the edge of ICP1 but still fits because it automates support operations and user-facing service flows.
- **Public signal:** Public positioning emphasizes AI support automation and faster resolution.
- **Likely buyer:** Head of Support, VP Product
- **Difficulty:** Access: Medium-Low; Onboarding: Low-Medium
- **Why difficulty is at that level:** Smaller and more approachable than incumbents, with a narrower workflow surface that could make for a faster pilot.

Regal.ai

- **HQ / ICP:** New York, USA / ICP1
- **Why it matches:** Regal.ai runs customer lifecycle and contact-center workflows where agentic execution quality affects revenue and service performance.
- **Public signal:** Public positioning emphasizes AI phone agents and orchestrated customer journeys.
- **Likely buyer:** VP Contact Center, VP Growth Operations
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** The workflows are high value and concrete, though integrations still matter.

Dialpad

- **HQ / ICP:** San Francisco, USA / ICP1
- **Why it matches:** Dialpad sits in AI contact center and communications workflows where real-time decision quality matters.
- **Public signal:** Public positioning includes AI agents, voice automation, and real-time assistance.
- **Likely buyer:** VP CX Technology, Head of AI Product
- **Difficulty:** Access: Medium-High; Onboarding: Medium
- **Why difficulty is at that level:** The fit is real, but buyers may initially anchor on productivity instead of execution control.

Five9

- **HQ / ICP:** San Ramon, USA / ICP1
- **Why it matches:** Five9 is a strong category fit through customer experience automation and AI-driven contact-center workflows.
- **Public signal:** Public positioning emphasizes AI, automation, and customer experience workflows.
- **Likely buyer:** VP Contact Center Technology, CTO
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** A strong target but not an easy first one due to enterprise scale and incumbent complexity.

Kognitos

- **HQ / ICP:** San Jose, USA / ICP1
- **Why it matches:** Kognitos is close to the wedge because it turns natural language operating procedures into automation.
- **Public signal:** Public positioning centers on AI-powered automation for enterprise operations.
- **Likely buyer:** VP Automation, CIO, Head of Platform
- **Difficulty:** Access: Medium; Onboarding: Medium-High
- **Why difficulty is at that level:** The conceptual fit is excellent, but pilots still require process mapping and operational data access.

ICP2: Experimentation Labs — 30 companies

Anthropic

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Anthropic is a high-end ICP2 target because it is deeply involved in agentic behavior, model safety, evals, and deployment quality.
- **Public signal:** Public launches around tool use, agents, and model safety show ongoing experimentation with runtime behavior.
- **Likely buyer:** Head of Research Platform, Head of Applied AI, VP Engineering
- **Difficulty:** Access: Very High; Onboarding: High
- **Why difficulty is at that level:** The fit is conceptually excellent, but access is difficult and the internal team is extremely sophisticated.

OpenAI

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** OpenAI is an experimentation-heavy environment where agent reliability, tool execution, and evaluation are active frontier concerns.
- **Public signal:** Public materials on agents, structured outputs, evals, and customer stories point to deep runtime behavior work.
- **Likely buyer:** Head of Applied Research, Product Engineering Lead, Platform Lead
- **Difficulty:** Access: Very High; Onboarding: High
- **Why difficulty is at that level:** A strategic dream target, but extremely hard to access and unlikely to move quickly on an external wedge.

Cohere

- **HQ / ICP:** Toronto, Canada / ICP2
- **Why it matches:** Cohere is a strong Canadian ICP2 target because it is deeply enterprise-facing while still experimentation-heavy.
- **Public signal:** Public positioning includes enterprise AI, retrieval, agents, and deployment tooling.
- **Likely buyer:** VP Product, Head of Platform AI, CTO
- **Difficulty:** Access: High; Onboarding: Medium-High
- **Why difficulty is at that level:** Reachable relative to frontier giants, but still technically advanced and partner selective.

Character.AI

- **HQ / ICP:** Menlo Park, USA / ICP2
- **Why it matches:** Character.AI is experimentation-heavy by default because product quality depends on live agent behavior and iteration.
- **Public signal:** Public product direction and company positioning center on rich AI characters and agent-like interactions at scale.
- **Likely buyer:** Head of Product, Head of Platform, CTO
- **Difficulty:** Access: High; Onboarding: Medium
- **Why difficulty is at that level:** The runtime behavior challenge is real, but the external design-partner angle may need careful framing.

Cognition

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Cognition is building autonomous software agents, which is one of the purest execution-control problem spaces.
- **Public signal:** Public product messaging around Devin and software agents shows heavy live experimentation.
- **Likely buyer:** Head of Agent Infrastructure, CTO, VP Engineering
- **Difficulty:** Access: High; Onboarding: Medium-High
- **Why difficulty is at that level:** The fit is strong, but the buyer bar will be very technical and skeptical.

Adept

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Adept is a direct ICP2 match because it focuses on agents taking action in software environments.
- **Public signal:** Public positioning centers on enterprise agents that use software tools and workflows.
- **Likely buyer:** Head of Platform, CTO, VP Engineering
- **Difficulty:** Access: Medium-High; Onboarding: Medium-High
- **Why difficulty is at that level:** The fit is excellent, but implementation likely touches sensitive action surfaces and tool permissions.

Perplexity

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Perplexity is moving from answer engine toward more agentic behavior, which creates new runtime quality constraints.
- **Public signal:** Public product direction increasingly blends search, actions, and tool use.
- **Likely buyer:** Head of Product, CTO, Head of Applied AI
- **Difficulty:** Access: High; Onboarding: Medium
- **Why difficulty is at that level:** A high-profile target with strong fit, but less obvious as a first design partner than workflow-native startups.

Replit

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Replit is a strong ICP2 candidate because coding agents and execution loops are central to the product direction.
- **Public signal:** Public product launches around agents and coding automation show active experimentation with execution quality.
- **Likely buyer:** Head of AI, VP Engineering, CTO
- **Difficulty:** Access: Medium-High; Onboarding: Medium
- **Why difficulty is at that level:** Reachable relative to frontier labs, with a real need for runtime reliability in agent execution.

Anysphere (Cursor)

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Cursor is operating in a fast-moving agentic coding environment where reliability and tool behavior change quickly.
- **Public signal:** Public positioning around AI pair programming and agent-like coding assistance shows high iteration velocity.
- **Likely buyer:** Head of Product, CTO, Head of Agent Systems
- **Difficulty:** Access: High; Onboarding: Medium
- **Why difficulty is at that level:** The company is hot and in demand, but the technical fit is strong enough to justify targeted outreach.

Codeium (Windsurf)

- **HQ / ICP:** Mountain View, USA / ICP2
- **Why it matches:** Windsurf is a strong ICP2 target because coding agents operate across files, tools, and developer workflows.
- **Public signal:** Public product direction highlights agentic coding and workflow automation inside the IDE.
- **Likely buyer:** Head of Product, VP Engineering, CTO
- **Difficulty:** Access: Medium-High; Onboarding: Medium
- **Why difficulty is at that level:** The fit is real and onboarding can be bounded to one agent path if the right champion exists.

Augment Code

- **HQ / ICP:** USA / ICP2
- **Why it matches:** Augment Code fits ICP2 because it lives in code generation, agent loops, and developer workflow experimentation.
- **Public signal:** Public positioning emphasizes coding agents and AI-native software development assistance.
- **Likely buyer:** VP Engineering, Head of Product, CTO
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** Smaller and more focused than top frontier labs, making it more plausible for a real design-partner motion.

Poolside

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Poolside is a serious ICP2 candidate because it is building coding intelligence where evaluation and runtime confidence are central.
- **Public signal:** Public positioning ties the company to software development agents and high-performance code generation.
- **Likely buyer:** Head of Research Engineering, CTO
- **Difficulty:** Access: High; Onboarding: Medium
- **Why difficulty is at that level:** The company is advanced, but still in a domain where execution-layer control is highly relevant.

Magic

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Magic fits because it is pushing autonomous coding systems and needs reliable execution against complex workflows.
- **Public signal:** Public positioning emphasizes software engineering automation and advanced AI systems.
- **Likely buyer:** Head of Engineering, CTO, Research Lead
- **Difficulty:** Access: High; Onboarding: Medium
- **Why difficulty is at that level:** The target is strategically relevant but technically demanding.

LangChain

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** LangChain is almost a textbook ICP2 candidate because the company lives in agents, orchestration, and workflow behavior.
- **Public signal:** Public product lines around LangGraph, LangSmith, and agent frameworks show direct proximity to runtime instability.
- **Likely buyer:** VP Product, CTO, Head of Platform
- **Difficulty:** Access: Medium; Onboarding: Low-Medium
- **Why difficulty is at that level:** The company deeply understands the problem, which raises the bar but also means messaging can be precise.

LlamaIndex

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** LlamaIndex is close to the core of agent data flows, orchestration, and experimentation.
- **Public signal:** Public positioning focuses on agents, data frameworks, and production AI workflows.
- **Likely buyer:** Head of Product, CTO, VP Engineering
- **Difficulty:** Access: Medium; Onboarding: Low-Medium
- **Why difficulty is at that level:** A good target because the company is technical, relevant, and not as inaccessible as the biggest labs.

Braintrust

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Braintrust is a very good fit because it already lives in evals and quality workflows, but does not itself close the execution-control gap.
- **Public signal:** Public positioning centers on AI evals, testing, and iteration loops.
- **Likely buyer:** CTO, Head of Product, Head of AI Quality
- **Difficulty:** Access: Medium; Onboarding: Low
- **Why difficulty is at that level:** Buyer education is easier because the team already thinks in quality and feedback loops.

Arize AI

- **HQ / ICP:** Berkeley, USA / ICP2
- **Why it matches:** Arize is strongly aligned because it sits in observability and eval infrastructure where customers still need runtime enforcement elsewhere.
- **Public signal:** Public positioning emphasizes observability, evals, and production iteration at scale.

- **Likely buyer:** VP Product, CTO, Head of AI Platform
- **Difficulty:** Access: Medium; Onboarding: Low-Medium
- **Why difficulty is at that level:** The conceptual bridge from observability to execution control is very clear, though differentiation must be sharp.

Galileo

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Galileo fits for the same reason as Arize: it sees the observability and eval layer but not the enforcement layer.
- **Public signal:** Public messaging emphasizes end-to-end visibility and agent evaluation metrics.
- **Likely buyer:** VP Product, CTO, Head of AI Systems
- **Difficulty:** Access: Medium; Onboarding: Low-Medium
- **Why difficulty is at that level:** The team will quickly understand the category distinction if the wedge is explained well.

Maxim AI

- **HQ / ICP:** USA / ICP2
- **Why it matches:** Maxim AI is close to the reliability problem because it tracks how agents reason and progress across multi-step tasks.
- **Public signal:** Public positioning focuses on AI agent observability and evaluation across autonomous systems.
- **Likely buyer:** Head of Product, CTO, VP Engineering
- **Difficulty:** Access: Medium; Onboarding: Low
- **Why difficulty is at that level:** Good fit and easier conceptual onboarding because the pain is already in scope.

Patronus AI

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Patronus AI is a fit because it works on evaluation, safety, and trust in production AI systems.
- **Public signal:** Public positioning emphasizes evaluation and trust for LLM applications and agents.
- **Likely buyer:** Head of Product, CTO, Head of Safety Engineering
- **Difficulty:** Access: Medium; Onboarding: Low-Medium
- **Why difficulty is at that level:** The fit is intellectually strong, though commercial urgency depends on their roadmap.

Confident AI

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Confident AI directly serves the LLM evaluation workflow, making it close to teams that know quality gaps firsthand.
- **Public signal:** Public positioning emphasizes deterministic metrics, tailored validation, and pre/post deployment evals.
- **Likely buyer:** CTO, Head of Product, Founder
- **Difficulty:** Access: Medium-Low; Onboarding: Low
- **Why difficulty is at that level:** A smaller and more focused target with category fluency, making it a plausible design-partner candidate.

Weights & Biases

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Weights & Biases is relevant because it sits in experiment tracking and model iteration where production-agent behavior is a growing frontier.
- **Public signal:** Public product direction includes support for AI development, experimentation, and monitoring workflows.
- **Likely buyer:** VP Product, Head of MLOps, CTO
- **Difficulty:** Access: High; Onboarding: Medium
- **Why difficulty is at that level:** Highly credible target, but broad product scope and existing internal sophistication raise the bar.

Vellum

- **HQ / ICP:** New York, USA / ICP2
- **Why it matches:** Vellum is a strong ICP2 fit because it is workflow-first around prompts, evaluation, and deployment.
- **Public signal:** Public positioning emphasizes prompt orchestration, evaluation, and enterprise AI automation.
- **Likely buyer:** Founder, Head of Product, CTO
- **Difficulty:** Access: Medium-Low; Onboarding: Low
- **Why difficulty is at that level:** A focused startup in the right part of the stack and likely more open to design-partner collaboration.

Lamini

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Lamini fits because enterprise model deployment and quality control are core to its value proposition.
- **Public signal:** Public messaging focuses on enterprise AI deployment and customization at scale.
- **Likely buyer:** CTO, VP Engineering, Head of Platform
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** A fit if the wedge is framed as runtime reliability for applied systems rather than model training alone.

Hebbia

- **HQ / ICP:** New York, USA / ICP2
- **Why it matches:** Hebbia is a fit because it uses agentic workflows for knowledge work where reliability and repeatability matter.
- **Public signal:** Public positioning centers on AI agents for high-value professional work.
- **Likely buyer:** VP Product, CTO, Head of Applied AI
- **Difficulty:** Access: Medium-High; Onboarding: Medium
- **Why difficulty is at that level:** The buyer case is strong, but access and differentiation may be harder due to category heat.

Stack AI

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Stack AI is an experimentation-heavy workflow company where building and deploying AI workflows is the product itself.
- **Public signal:** Public positioning emphasizes no-code AI workflows and enterprise automation.
- **Likely buyer:** Founder, CTO, Head of Product
- **Difficulty:** Access: Medium-Low; Onboarding: Low
- **Why difficulty is at that level:** A good design-partner shape because the workflow surface is direct and the team is likely still iterating fast.

Factory

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Factory fits because it is building software engineering AI systems where execution reliability is a real problem.
- **Public signal:** Public product direction centers on agentic software development and developer workflow acceleration.
- **Likely buyer:** CTO, Head of Engineering, Product Lead
- **Difficulty:** Access: Medium-High; Onboarding: Medium
- **Why difficulty is at that level:** The fit is meaningful, but the technical bar will be high.

Gantry

- **HQ / ICP:** USA / ICP2
- **Why it matches:** Gantry is relevant because it operates near model evaluation, production quality, and AI reliability workflows.
- **Public signal:** Public positioning centers on improving and governing production AI systems.
- **Likely buyer:** Founder, CTO, Head of ML Platform
- **Difficulty:** Access: Medium-Low; Onboarding: Low
- **Why difficulty is at that level:** A good target for category conversation because the team likely understands quality and governance tradeoffs well.

All Hands AI

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** All Hands AI fits because open agent systems and developer workflows expose real execution instability.
- **Public signal:** Public positioning emphasizes agentic software development and open-ended execution.
- **Likely buyer:** Founder, CTO, Head of Agent Systems
- **Difficulty:** Access: Medium-Low; Onboarding: Low-Medium
- **Why difficulty is at that level:** Smaller and more experimentation-heavy, which helps access, though budgets may be smaller than enterprise targets.

Writer

- **HQ / ICP:** San Francisco, USA / ICP2
- **Why it matches:** Writer is a strong fit because it sits between enterprise deployment, evaluation, governance, and workflow AI.
- **Public signal:** Public positioning includes AI apps, agents, governance, and enterprise control.
- **Likely buyer:** GM Platform, CTO, VP Product
- **Difficulty:** Access: Medium-High; Onboarding: Medium-High
- **Why difficulty is at that level:** Strategically aligned, but already mature enough to expect clear proof and tight positioning.

ICP3: Regulated Systems — 20 companies

Abridge

- **HQ / ICP:** Pittsburgh, USA / ICP3
- **Why it matches:** Abridge is a prime ICP3 target because it automates clinical workflows where traceability and correctness matter deeply.
- **Public signal:** Public positioning centers on ambient clinical documentation and health system deployments.
- **Likely buyer:** CMIO, VP Engineering, CTO
- **Difficulty:** Access: High; Onboarding: Very High
- **Why difficulty is at that level:** Healthcare security, EHR integration, HIPAA, and clinical change management make pilots heavy but valuable.

Ambience Healthcare

- **HQ / ICP:** San Francisco, USA / ICP3
- **Why it matches:** Ambience Healthcare fits because it automates documentation and coding workflows in regulated clinical settings.
- **Public signal:** Public positioning emphasizes clinical AI, documentation, and workflow efficiency for providers.
- **Likely buyer:** CMIO, VP Product, CTO
- **Difficulty:** Access: High; Onboarding: Very High
- **Why difficulty is at that level:** The buyer case is strong, but implementation touches clinical systems and compliance from day one.

Hippocratic AI

- **HQ / ICP:** Palo Alto, USA / ICP3
- **Why it matches:** Hippocratic AI is directly in regulated healthcare agent behavior, where runtime safety and explainability are core.
- **Public signal:** Public positioning emphasizes healthcare AI agents and safety-focused deployment.
- **Likely buyer:** Chief Medical Officer, CTO, Head of Safety
- **Difficulty:** Access: High; Onboarding: Very High
- **Why difficulty is at that level:** A strong category fit, but oversight and regulatory scrutiny make pilots complex.

Suki

- **HQ / ICP:** Redwood City, USA / ICP3
- **Why it matches:** Suki automates clinician workflows and documentation, making execution quality and auditability highly relevant.
- **Public signal:** Public product messaging focuses on AI assistants for clinicians and workflow efficiency.
- **Likely buyer:** CMIO, VP Engineering, CTO
- **Difficulty:** Access: High; Onboarding: Very High
- **Why difficulty is at that level:** Clinical integration and physician workflow sensitivity raise the difficulty materially.

Qventus

- **HQ / ICP:** Mountain View, USA / ICP3
- **Why it matches:** Qventus fits because it uses AI to improve hospital operations and care workflows, which are high-stakes execution environments.
- **Public signal:** Public positioning emphasizes healthcare operations automation and workflow optimization.
- **Likely buyer:** COO, VP Digital Health, CTO
- **Difficulty:** Access: High; Onboarding: Very High
- **Why difficulty is at that level:** Operational value is large, but hospital systems and process ownership make onboarding slow.

Commure

- **HQ / ICP:** Mountain View, USA / ICP3
- **Why it matches:** Commure is compelling because it sits on top of healthcare operational infrastructure and AI workflows.
- **Public signal:** Public positioning blends healthcare software infrastructure with AI-enabled operational workflows.
- **Likely buyer:** CTO, VP Platform, CIO Office
- **Difficulty:** Access: High; Onboarding: Very High
- **Why difficulty is at that level:** Strategic target with real fit, but technical and compliance scope are both heavy.

Infinitus Systems

- **HQ / ICP:** San Francisco, USA / ICP3
- **Why it matches:** Infinitus is a strong fit because AI voice agents in healthcare payer-provider workflows need control, auditability, and safe execution.
- **Public signal:** Public positioning emphasizes automated healthcare phone workflows and operational efficiency.
- **Likely buyer:** VP Engineering, VP Operations, CTO
- **Difficulty:** Access: Medium-High; Onboarding: Very High
- **Why difficulty is at that level:** The workflow value is high, but regulated telephony and healthcare process integration raise complexity.

Augmedix

- **HQ / ICP:** San Francisco, USA / ICP3
- **Why it matches:** Augmedix fits because it automates clinical documentation and workflow tasks in a regulated care environment.
- **Public signal:** Public positioning emphasizes ambient documentation and clinical workflow support.
- **Likely buyer:** CMIO, CTO, VP Product
- **Difficulty:** Access: High; Onboarding: Very High
- **Why difficulty is at that level:** Healthcare deployment complexity is unavoidable even if the use case is clear.

Sardine

- **HQ / ICP:** San Francisco, USA / ICP3
- **Why it matches:** Sardine is a very good fintech ICP3 target because fraud and risk workflows need runtime enforcement, not just visibility.
- **Public signal:** Public positioning centers on fraud, compliance, and risk decisioning automation.
- **Likely buyer:** Head of Risk Platform, CTO, VP Engineering
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** The fit is strong, but integration into sensitive risk controls and data flows raises the implementation bar.

Alloy

- **HQ / ICP:** New York, USA / ICP3
- **Why it matches:** Alloy automates identity and risk decisions, which makes explainability and policy-safe execution crucial.
- **Public signal:** Public positioning emphasizes identity decisioning, fraud prevention, and compliance workflows.
- **Likely buyer:** Head of Risk Engineering, CTO, VP Product
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** A good wedge exists, but buyer scrutiny will be high because risk actions touch regulated decisions.

Ocrolus

- **HQ / ICP:** New York, USA / ICP3
- **Why it matches:** Ocrolus fits because it automates document-heavy underwriting and financial review workflows where execution errors matter.
- **Public signal:** Public positioning focuses on document automation for lending and financial decision workflows.
- **Likely buyer:** VP Engineering, Head of Risk Operations, CTO
- **Difficulty:** Access: Medium; Onboarding: High
- **Why difficulty is at that level:** The use case is concrete, but onboarding still requires access to sensitive financial data and decision rules.

Trullion

- **HQ / ICP:** New York, USA / ICP3
- **Why it matches:** Trullion is a fit because it automates accounting and audit workflows where traceability and correctness are core.
- **Public signal:** Public positioning emphasizes accounting automation, audit support, and structured financial workflows.
- **Likely buyer:** CTO, Head of Product, VP Finance Systems
- **Difficulty:** Access: Medium; Onboarding: Medium-High
- **Why difficulty is at that level:** The workflow scope can be bounded, but financial-data sensitivity and process ownership increase complexity.

AppZen

- **HQ / ICP:** San Jose, USA / ICP3
- **Why it matches:** AppZen is directly in finance audit and compliance automation, which is close to the wedge.
- **Public signal:** Public positioning centers on finance AI for auditing, expense review, and compliance workflows.
- **Likely buyer:** VP Finance Systems, CTO, Head of AI Platform
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** The fit is good, but implementation touches financial controls and approval chains.

Ramp

- **HQ / ICP:** New York, USA / ICP3
- **Why it matches:** Ramp is a valid ICP3 target because it runs AI and automation inside finance operations where mistakes affect spend and controls.
- **Public signal:** Public product direction increasingly includes AI automation across finance workflows.
- **Likely buyer:** Head of AI, CTO, VP Finance Product
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** The company is highly sophisticated and fast-moving, but the workflow-control problem is real.

Harvey

- **HQ / ICP:** San Francisco, USA / ICP3
- **Why it matches:** Harvey is one of the best legal ICP3 targets because it already runs workflow agents for law firms and corporate legal teams.
- **Public signal:** Public positioning explicitly mentions workflow agents and legal task automation across compliance, contracts, and research.
- **Likely buyer:** Head of Product, CTO, Head of Applied AI
- **Difficulty:** Access: High; Onboarding: High
- **Why difficulty is at that level:** The fit is excellent, but legal buyers expect trust, rigor, and very specific proof.

Clio

- **HQ / ICP:** Vancouver, Canada / ICP3
- **Why it matches:** Clio is a strong ICP3 target because legal practice workflows are structured, high-value, and regulation-sensitive.
- **Public signal:** Public product direction increasingly includes AI capabilities for legal work and operations.
- **Likely buyer:** CTO, VP Product, Head of AI
- **Difficulty:** Access: Medium-High; Onboarding: Medium-High
- **Why difficulty is at that level:** Good category fit with reachable buyers, though workflow credibility matters a lot.

Spellbook

- **HQ / ICP:** Toronto, Canada / ICP3
- **Why it matches:** Spellbook is a close match because contract review and drafting workflows require reliable structured behavior.
- **Public signal:** Public positioning centers on AI for contract drafting and legal review inside existing legal workflows.
- **Likely buyer:** Founder, CTO, VP Product
- **Difficulty:** Access: Medium; Onboarding: Medium
- **Why difficulty is at that level:** A focused legal AI company with a bounded workflow surface and realistic design-partner shape.

Streamline AI

- **HQ / ICP:** San Francisco, USA / ICP3
- **Why it matches:** Streamline AI is directly in legal intake, triage, and workflow automation, making execution control highly relevant.
- **Public signal:** Public positioning emphasizes intelligent intake, triage, and legal workflow automation.
- **Likely buyer:** Head of Legal Ops Product, CTO, Founder
- **Difficulty:** Access: Medium; Onboarding: High
- **Why difficulty is at that level:** The buyer pain is clear, but implementation touches legal operations, approvals, and compliance boundaries.

Evisort

- **HQ / ICP:** San Mateo, USA / ICP3
- **Why it matches:** Evisort fits because contract and legal workflow automation depend on repeatable, explainable execution.
- **Public signal:** Public product direction centers on contract AI and workflow automation for legal teams.
- **Likely buyer:** VP Product, CTO, Head of Legal AI
- **Difficulty:** Access: Medium; Onboarding: High
- **Why difficulty is at that level:** The fit is real, though legal stakeholders will demand clear evidence of control and traceability.

EvenUp

- **HQ / ICP:** San Francisco, USA / ICP3
- **Why it matches:** EvenUp is a fit because it automates parts of legal claims workflows where reliability and auditability matter.
- **Public signal:** Public positioning emphasizes AI for injury law workflows and case preparation.
- **Likely buyer:** CTO, VP Product, Head of Operations
- **Difficulty:** Access: Medium-High; Onboarding: High
- **Why difficulty is at that level:** The domain is high value, but trust and workflow correctness requirements will be strict.

Highest-priority starting targets

If the goal is to maximize the chance of landing a real design partner quickly, these are the best first outreach targets from the full 100-company universe:

- Ada
- Decagon
- Forethought
- Observe.AI
- Gorgias
- Intercom
- Assembled
- Stack AI
- Vellum
- Confident AI
- Spellbook
- Streamline AI
- Sardine
- Alloy
- Infinitus Systems

These companies combine three useful properties: visible workflow automation, public evidence of agent or orchestration complexity, and a better chance of actually taking a design-partner conversation than the biggest incumbents.

Suggested next operating pass

The highest-leverage next step would be to turn this target universe into a narrower outreach file with:

- named likely buyer titles for each account,
- best entry wedge for the first message,
- one-line proof point to reference in outreach,
- and a short ranking of who to contact first this month.

That would convert this research document into an actual founder outbound asset.